

```

from __future__ import print_function

import numpy as np
from hmwk5_2.classifiers.linear_svm import *
from hmwk5_2.classifiers.softmax import *

class LinearClassifier(object):

    def __init__(self):
        self.W = None

    def train(self, X, y, learning_rate=1e-3, reg=1e-5, num_iters=100,
              batch_size=200, verbose=False):
        """
        Train this linear classifier using stochastic gradient descent.

        Inputs:
        - X: A numpy array of shape (N, D) containing training data; there are N
            training samples each of dimension D.
        - y: A numpy array of shape (N,) containing training labels; y[i] = c
            means that X[i] has label 0 ≤ c < C for C classes.
        - learning_rate: (float) learning rate for optimization.
        - reg: (float) regularization strength.
        - num_iters: (integer) number of steps to take when optimizing
        - batch_size: (integer) number of training examples to use at each step.
        - verbose: (boolean) If true, print progress during optimization.

        Outputs:
        A list containing the value of the loss function at each training iteration.
        """
        num_train, dim = X.shape
        num_classes = np.max(y) + 1 # assume y takes values 0...K-1 where K is number of classes
        if self.W is None:
            # lazily initialize W
            self.W = 0.001 * np.random.randn(dim, num_classes)

        # Run stochastic gradient descent to optimize W
        loss_history = []
        for it in range(num_iters):
            X_batch = None
            y_batch = None

            #####
            # TODO:
            # Sample batch_size elements from the training data and their
            # corresponding labels to use in this round of gradient descent.
            # Store the data in X_batch and their corresponding labels in
            # y_batch; after sampling X_batch should have shape (dim, batch_size)
            # and y_batch should have shape (batch_size,)
            #
            # Hint: Use np.random.choice to generate indices. Sampling with
            # replacement is faster than sampling without replacement.
            #####
            indxs = np.random.choice(num_train, batch_size)
            X_batch = X[indxs]
            y_batch = y[indxs]

            pass
            #####
            #
            # END OF YOUR CODE
            #####

            # evaluate loss and gradient
            loss, grad = self.loss(X_batch, y_batch, reg)

```

```
loss_history.append(loss)

# perform parameter update
#####
# TODO:
# Update the weights using the gradient and the learning rate.
#####
self.W -= learning_rate*grad
pass
#####
#                               END OF YOUR CODE
#####

if verbose and it % 100 == 0:
    print('iteration %d / %d: loss %f' % (it, num_iters, loss))

return loss_history

def predict(self, X):
    """
    Use the trained weights of this linear classifier to predict labels for
    data points.

    Inputs:
    - X: A numpy array of shape (N, D) containing training data; there are N
        training samples each of dimension D.

    Returns:
    - y_pred: Predicted labels for the data in X. y_pred is a 1-dimensional
        array of length N, and each element is an integer giving the predicted
        class.
    """
    y_pred = np.zeros(X.shape[0])
    #####
    # TODO:
    # Implement this method. Store the predicted labels in y_pred.
    #####
    y_pred = np.argmax(np.dot(X, self.W), axis = 1)
    pass
    #####
    #                               END OF YOUR CODE
    #####
    return y_pred

def loss(self, X_batch, y_batch, reg):
    """
    Compute the loss function and its derivative.
    Subclasses will override this.

    Inputs:
    - X_batch: A numpy array of shape (N, D) containing a minibatch of N
        data points; each point has dimension D.
    - y_batch: A numpy array of shape (N,) containing labels for the minibatch.
    - reg: (float) regularization strength.

    Returns: A tuple containing:
    - loss as a single float
    - gradient with respect to self.W; an array of the same shape as W
    """
    pass
```

```
class LinearSVM(LinearClassifier):
```

```
""" A subclass that uses the Multiclass SVM loss function """
```

```
def loss(self, X_batch, y_batch, reg):  
    return svm_loss_vectorized(self.W, X_batch, y_batch, reg)
```

```
class Softmax(LinearClassifier):
```

```
    """ A subclass that uses the Softmax + Cross-entropy loss function """
```

```
def loss(self, X_batch, y_batch, reg):  
    return softmax_loss_vectorized(self.W, X_batch, y_batch, reg)
```