

```
import numpy as np
from random import shuffle
```

```
def svm_loss_naive(W, X, y, reg):
```

```
    """
```

```
    Structured SVM loss function, naive implementation (with loops).
```

```
    Inputs have dimension D, there are C classes, and we operate on minibatches
    of N examples.
```

```
    Inputs:
```

- W: A numpy array of shape (D, C) containing weights.
- X: A numpy array of shape (N, D) containing a minibatch of data.
- y: A numpy array of shape (N,) containing training labels; $y[i] = c$ means that $X[i]$ has label c , where $0 \leq c < C$.
- reg: (float) regularization strength

```
    Returns a tuple of:
```

- loss as single float
- gradient with respect to weights W; an array of same shape as W

```
    """
```

```
    dW = np.zeros(W.shape) # initialize the gradient as zero
```

```
    # compute the loss and the gradient
```

```
    num_classes = W.shape[1]
```

```
    num_train = X.shape[0]
```

```
    loss = 0.0
```

```
    for i in range(num_train):
```

```
        scores = X[i].dot(W)
```

```
        correct_class_score = scores[y[i]]
```

```
        for j in range(num_classes):
```

```
            if j == y[i]:
```

```
                continue
```

```
            margin = scores[j] - correct_class_score + 1 # note delta = 1
```

```
            if margin > 0:
```

```
                loss += margin
```

```
                dW[:, y[i]] -= X[i, :]
```

```
                dW[:, j] += X[i, :]
```

```
    # Right now the loss is a sum over all training examples, but we want it
    # to be an average instead so we divide by num_train.
```

```
    loss /= num_train
```

```
    dW /= num_train
```

```
    # Add regularization to the loss.
```

```
    loss += reg * np.sum(W * W)
```

```
    # Add regularization to the derivative
```

```
    dW += reg * W
```

```
#####
# TODO:
# Compute the gradient of the loss function and store it dW.
# Rather than first computing the loss and then computing the derivative,
# it may be simpler to compute the derivative at the same time that the
# loss is being computed. As a result you may need to modify some of the
# code above to compute the gradient.
#####
```

```
return loss, dW
```

```

def svm_loss_vectorized(W, X, y, reg):
    """
    Structured SVM loss function, vectorized implementation.

    Inputs and outputs are the same as svm_loss_naive.
    """
    loss = 0.0
    dW = np.zeros(W.shape) # initialize the gradient as zero

    #####
    # TODO:
    # Implement a vectorized version of the structured SVM loss, storing the
    # result in loss.
    #####
    score = X.dot(W)

    yscore = score[np.arange(score.shape[0]),y]
    margin = np.maximum(0, score-np.matrix(yscore).T+1)
    margin[np.arange(X.shape[0]),y] = 0

    loss = np.sum(margin)

    # Average
    loss /= X.shape[0]

    # Add regularization
    loss += 0.5 * reg * np.sum(W * W)
    pass
    #####
    #
    # END OF YOUR CODE
    #
    #####

    #####
    # TODO:
    # Implement a vectorized version of the gradient for the structured SVM
    # loss, storing the result in dW.
    #
    # Hint: Instead of computing the gradient from scratch, it may be easier
    # to reuse some of the intermediate values that you used to compute the
    # loss.
    #####

    grad = np.zeros(score.shape)

    L = margin

    L[L > 0] = 1 #set the ones above 0 to 1 since they are support vectors that are contributing
    to gradient

    L[np.arange(score.shape[0]),y] -= np.sum(L, axis=1).T
    dW = np.dot(X.T,L)

    dW /= X.shape[0]
    dW += reg*W
    pass
    #####
    #
    # END OF YOUR CODE
    #
    #####

    return loss, dW

```